Class 7: Machine Learning

Nhi To

Today we wil explore unsupervised machine learning methods including clustering and dimensionality reduction methods.

Let's start by making up some data (where we know there are clear gorups) that we can use to test out different clustering methods.

We can use the 'rnorm()' function to help us here:

```
hist( rnorm(n=3000, mean=3) )
```

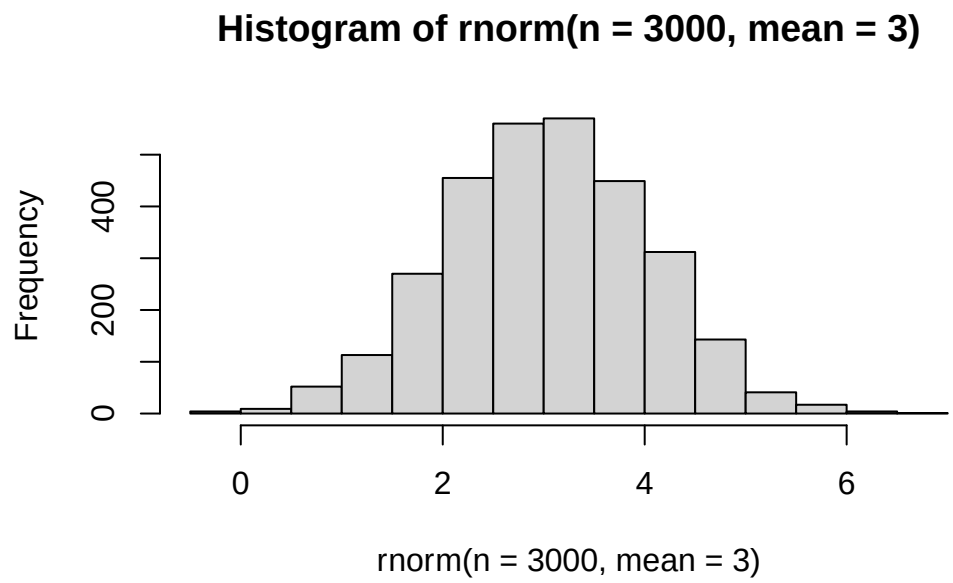

Make data with two "clusters'

```
x <- c( rnorm(30, mean=-3),
   rnorm(30, mean=+3) )

z <- cbind(x=x, rev(x))
head(z)
```

```
             x
[1,] -2.869380 3.291990
[2,] -2.510397 3.506521
[3,] -2.432219 2.553661
[4,] -1.262768 2.973889
[5,] -3.535593 3.270077
[6,] -2.839780 2.116586
```

```
plot(z)
```

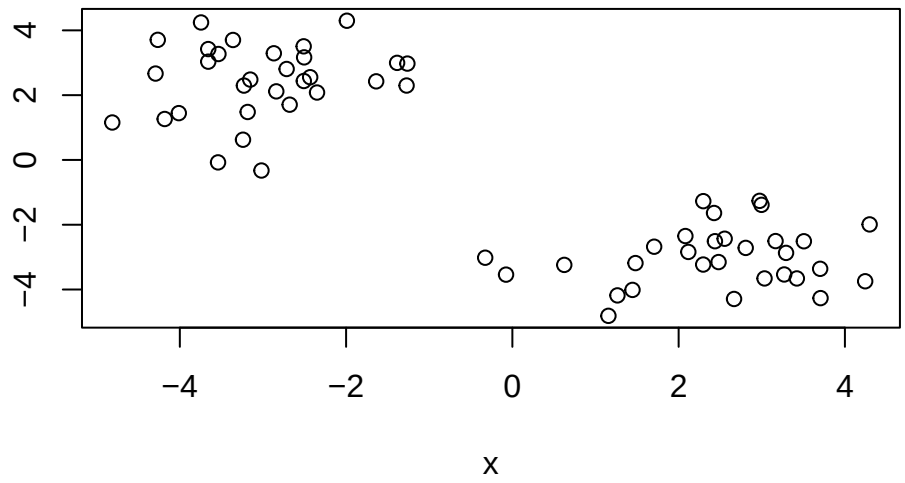

How big is 'z'

```
nrow(z)
```

```
[1] 60
```

```r
ncol(z)
```

```
[1] 2
```

K-means clustering

The main function in "base" R for K-means clustering is called 'kmeans()'

```r
k <- kmeans(z, centers= 2)
k
```

```
K-means clustering with 2 clusters of sizes 30, 30

Cluster means:
          x
1  2.434715 -2.994850
2 -2.994850  2.434715

Clustering vector:
 [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1

Within cluster sum of squares by cluster:
[1] 62.58684 62.58684
 (between_SS / total_SS =  87.6 %)

Available components:

[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```r
attributes(k)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"

$class
[1] "kmeans"
```

Q. How many points lie in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What component of our results tells us about the cluster membership (i.e. which pointlies in which cluster)?

```
k$cluster
```

```
 [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Q, Center of each cluster?

```
k$centers
```

```
          x
1  2.434715 -2.994850
2 -2.994850  2.434715
```

Q. Put this result info together and make a little "base R" plot of our clustering result. Also add the cluster center points to this plot.

```
plot (z, col= c("blue", "red"))
```

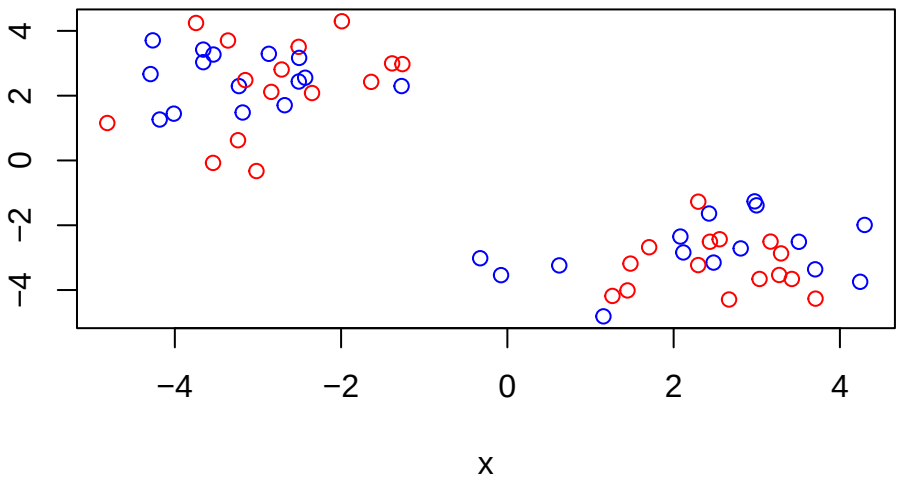

You can color by number

```
plot(z, col=c(1, 2))
```

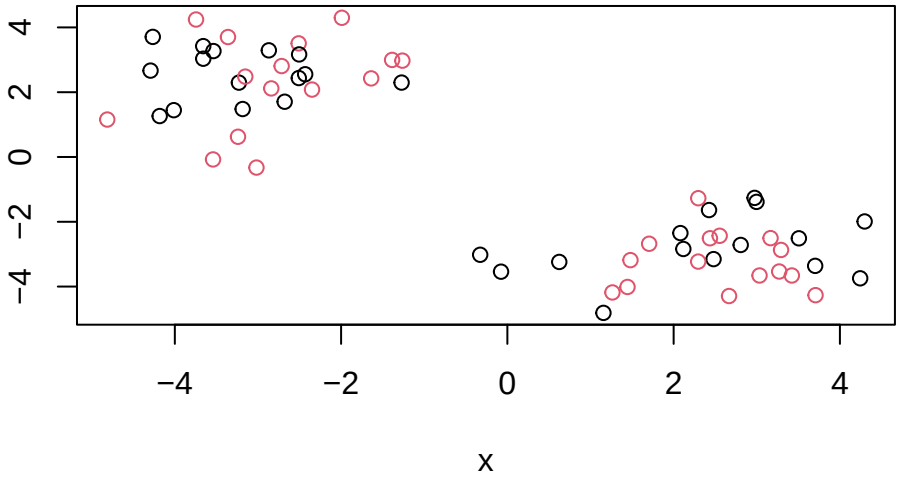

Plot colored by cluster membership:

```
plot(z, col= k$cluster)
points(k$centers, col="blue", pch=15)
```

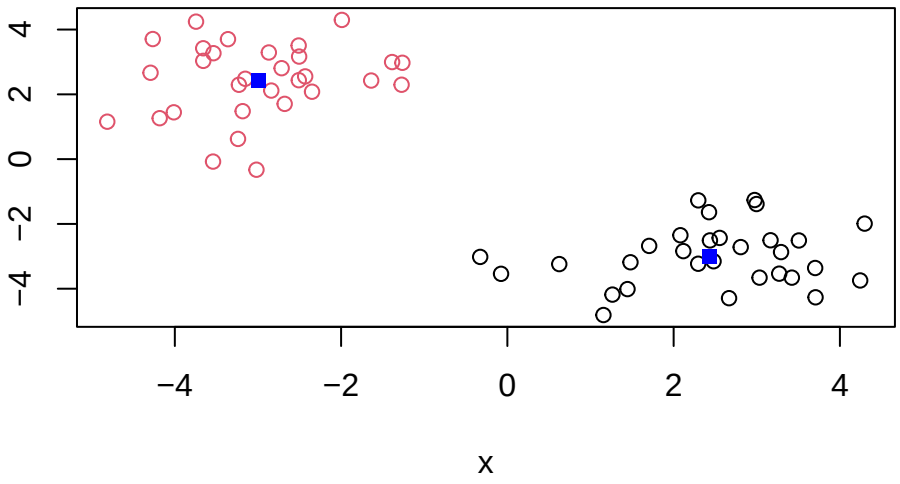

Q. Run kmeans on our input 'z' and define 4 clusters making the same result visualiation plot as plot as one (plot of z colored by cluster membership)

```
k4 <- kmeans(z, centers=4)
plot(z, col= k4$cluster)
```

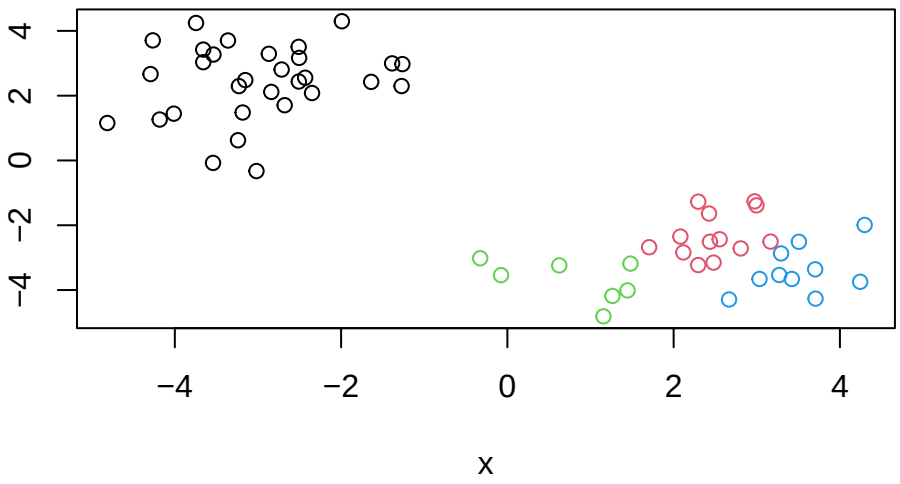

Hiearchial CLustering

The main function in base R for this called 'hclus()' it will take as input a distance matrix(key point is that you can't just give your "raw" data as input- you have to first calculate a distance martrix from your data).

```
d<- dist(z)
hc <- hclust(d)
hc
```

```
Call:
hclust(d = d)

Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

```
plot(hc)
abline(h=10, col="red")
```

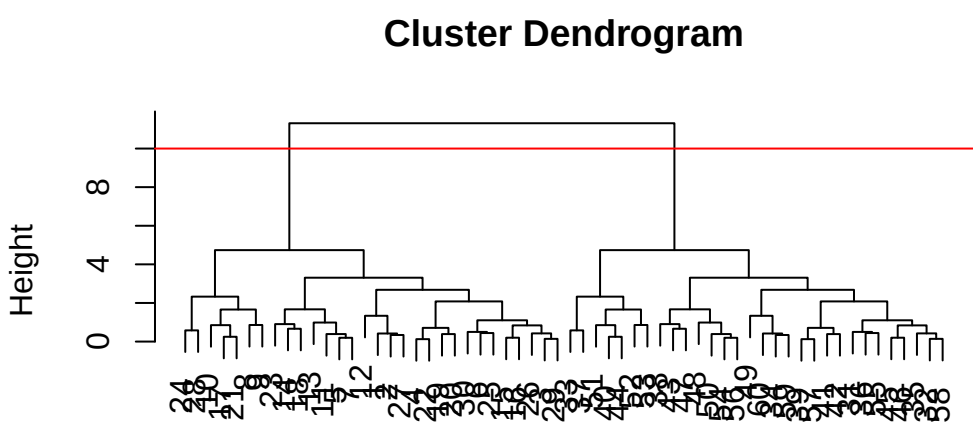

Once I inspect the "tree" I can "cut" the tree to yield my groupings or clusters. The function to do this is called 'cutree()'

```
grps<- cutree(hc, h=10)
```

```
plot(z, col=grps)
```

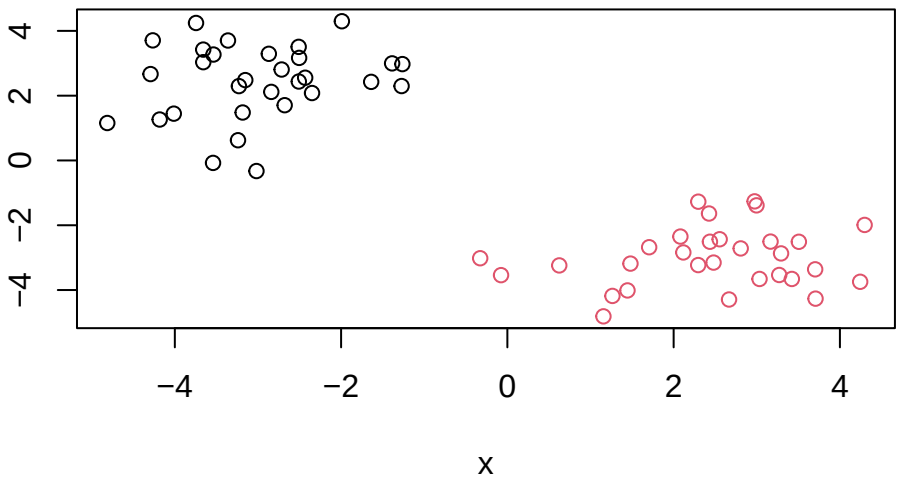

Hands on with Principal Component Analysis (PCA)

Let's examine some silly 17-dimensional data detailing food consumption in the UK (England, Scotland, Wales, and N. Ireland). Are these countries eating habits different or similr and if so how?

Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

```
                England Wales Scotland N.Ireland
Cheese              105   103      103        66
Carcass_meat        245   227      242       267
Other_meat          685   803      750       586
Fish                147   160      122        93
Fats_and_oils       193   235      184       209
Sugars              156   175      147       139
Fresh_potatoes      720   874      566      1033
```

```
Fresh_Veg                253   265      171       143
Other_Veg                488   570      418       355
Processed_potatoes       198   203      220       187
Processed_Veg            360   365      337       334
Fresh_fruit             1102  1137      957       674
Cereals                 1472  1582     1462      1494
Beverages                 57    73       53        47
Soft_drinks             1374  1256     1572      1506
Alcoholic_drinks         375   475      458       135
Confectionery             54    64       62        41
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

Answer: There are 17 rows and 4 columns

Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Answer:I approach "nrow()" and "ncol()" because it allows me to choose with one I would like to look at. If i use "dim()" I would be generating both.

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```

```
dim(x)
```

```
[1] 17  4
```

Q3. Changing what optional argument in the above barplot() function results in the following plot?

Answer:Changing the argument from barplot(as.matrix(x), beside=T, col=rainbow(nrow(x))) to barplot(as.matrix(x), beside=F, col=rainbow(nrow(x))) would result in a stacked barplot instead of a grouped barplot. The argument being changed was beside=T to beside=F.

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

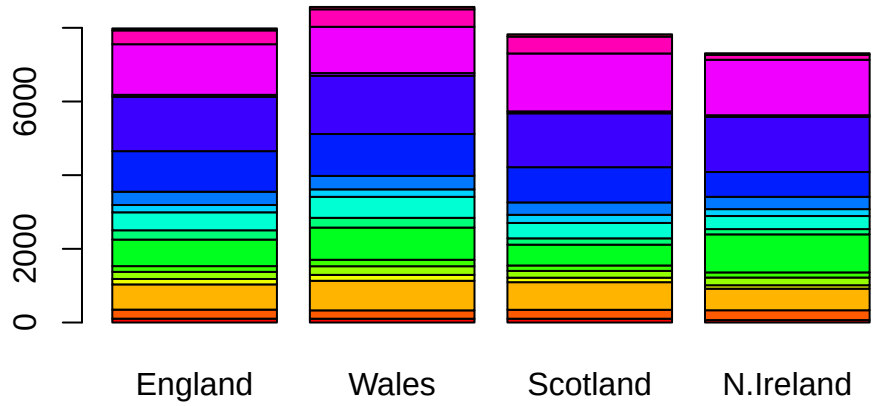

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Ans: if the points are diagonal on a given plot, it means that the values of the two compared countries are similar. If the point does not align diagonally, it means that they are not similar.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Ans: The main differences between N. Ireland the other countries of the UK are the consumption of potatoes, fresh fruit, soft drinks, and alcoholic drinks. Ireland has consumed significantly more fresh potatoes than other countries of the UK. Ireland has consumed significantly less fresh fruit than the other countries of the UK. Ireland has consumed much more soft drinks than Wales. Ireland has consumed significantly less alcoholic dirnks than the other countries of the UK.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```

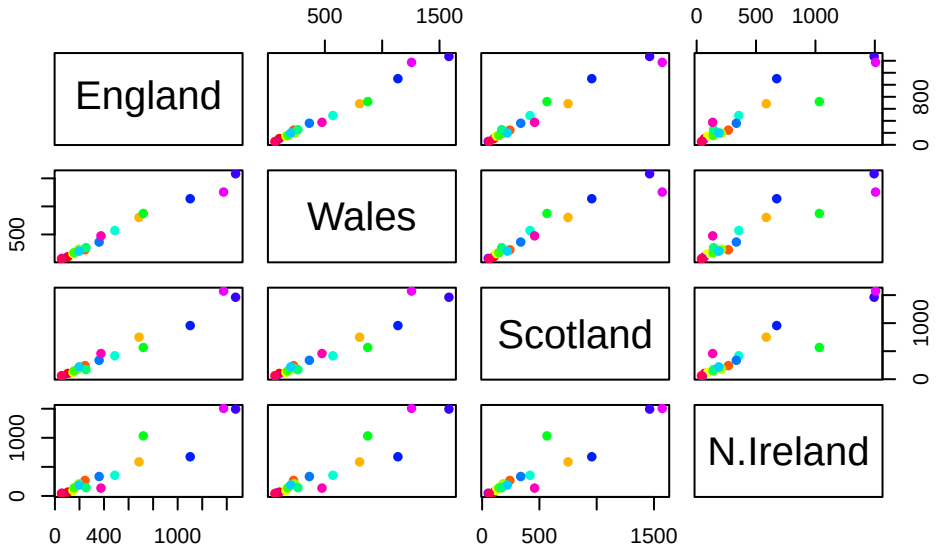

Looking at these types of "pairwise plots" can be helpful but it does not scale well and kind of sucks!There must be a better way…

PCA to the rescue!

The main function for PCA is base R is called 'prcomp()'. This function wants the transpose of our input data- i.e. the important foods in as columns and the countries as rows.

```
pca <- prcomp(t(x))
summary(pca)
```

```
Importance of components:
                            PC1      PC2      PC3        PC4
Standard deviation     324.1502 212.7478 73.87622 2.921e-14
Proportion of Variance   0.6744   0.2905  0.03503 0.000e+00
Cumulative Proportion    0.6744   0.9650  1.00000 1.000e+00
```

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```

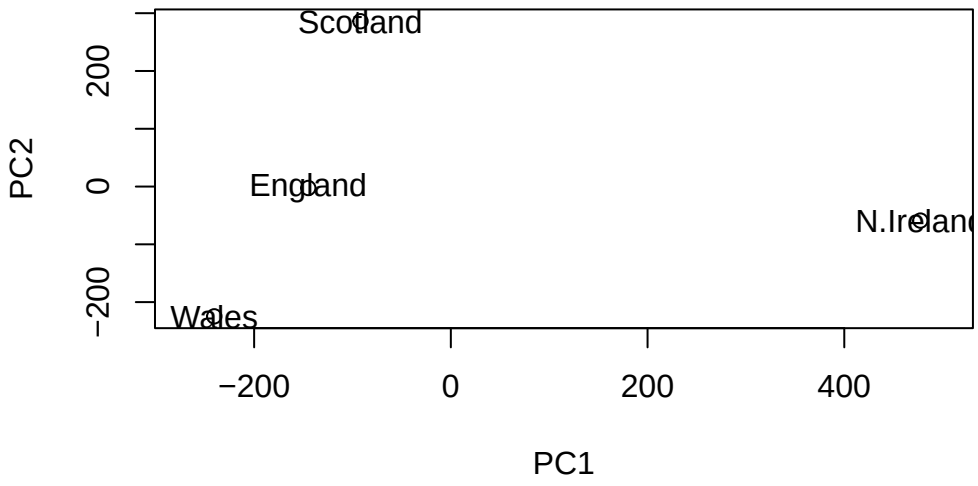

Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "darkgreen"))
```

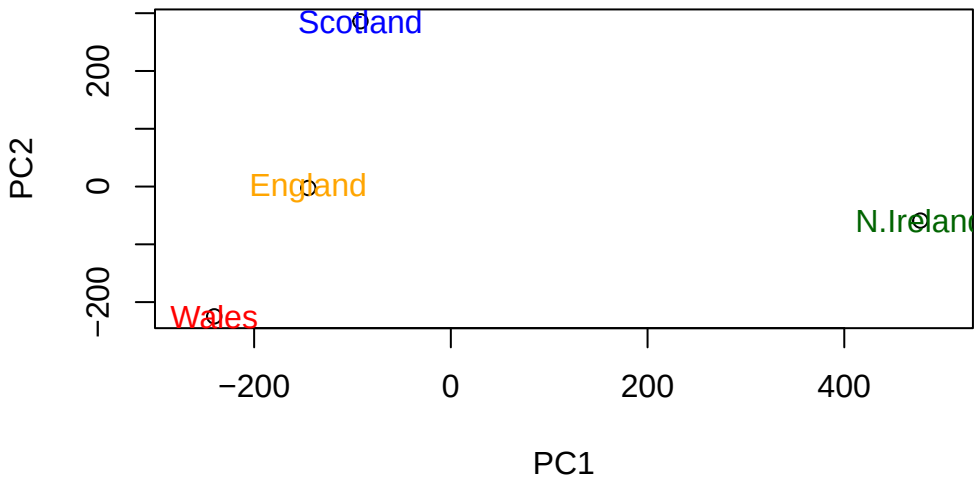

Looking at these types of "pairwise plots" can be helpful but it does not scale well and kind of sucks!There must be a better way…

Let's see what is in our PCA result object 'pca'

```
attributes(pca)
```

```
$names
[1] "sdev"     "rotation" "center"   "scale"    "x"

$class
[1] "prcomp"
```

The 'pca$x' result object is where we will focus first as this details how the countries are related to each other in terms of our new "axis" (a.k.a "PCs", "eigenvectors", etc.)

```
head(pca$x)
```

```
                PC1         PC2        PC3           PC4
England   -144.99315   -2.532999 105.768945 -9.152022e-15
Wales     -240.52915 -224.646925 -56.475555  5.560040e-13
Scotland   -91.86934  286.081786 -44.415495 -6.638419e-13
N.Ireland  477.39164  -58.901862  -4.877895  1.329771e-13
```

```
plot(pca$x[,1], pca$x[,2], pch=16, col=c("orange", "red", "blue", "darkgreen"))
```

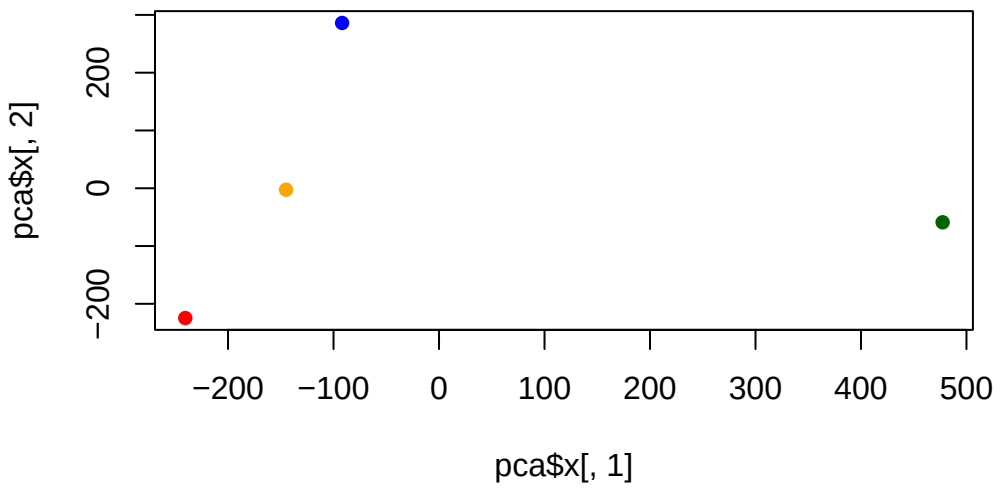

We can look at the so-called PC "loadings" result object to seehow the orignal foods contribute to our new PCs (i.e. how the original variables contribute to our new better PC variables)

```
pca$rotation[,1]
```

```
          Cheese       Carcass_meat         Other_meat               Fish
    -0.056955380         0.047927628       -0.258916658       -0.084414983
Fats_and_oils              Sugars     Fresh_potatoes          Fresh_Veg
    -0.005193623       -0.037620983        0.401402060       -0.151849942
   Other_Veg  Processed_potatoes      Processed_Veg         Fresh_fruit
    -0.243593729       -0.026886233       -0.036488269       -0.632640898
      Cereals           Beverages       Soft_drinks    Alcoholic_drinks
    -0.047702858       -0.026187756        0.232244140       -0.463968168
Confectionery
    -0.029650201
```

Calculating how much variation in the original data each PC accounts for

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29  4  0
```

```
z <- summary(pca)
z$importance
```

```
                            PC1       PC2      PC3          PC4
Standard deviation     324.15019 212.74780 73.87622 2.921348e-14
Proportion of Variance   0.67444   0.29052  0.03503 0.000000e+00
Cumulative Proportion    0.67444   0.96497  1.00000 1.000000e+00
```

Plot v

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```

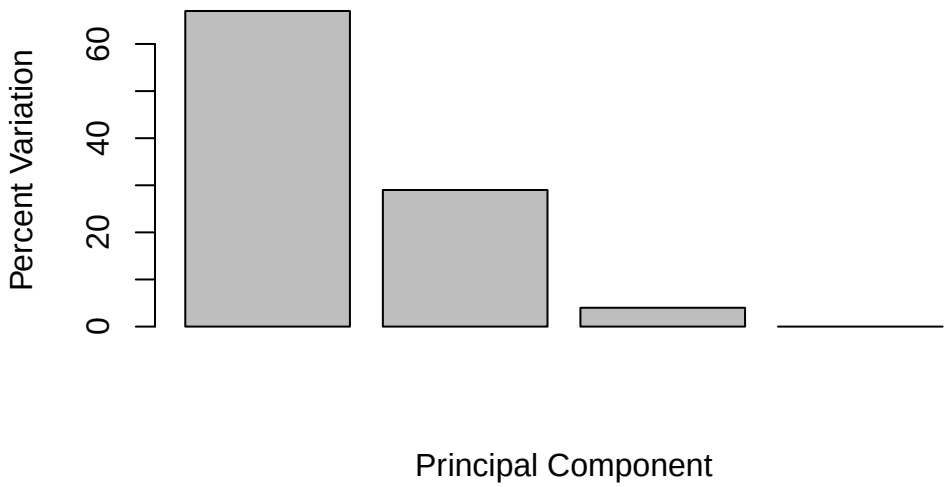


```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```

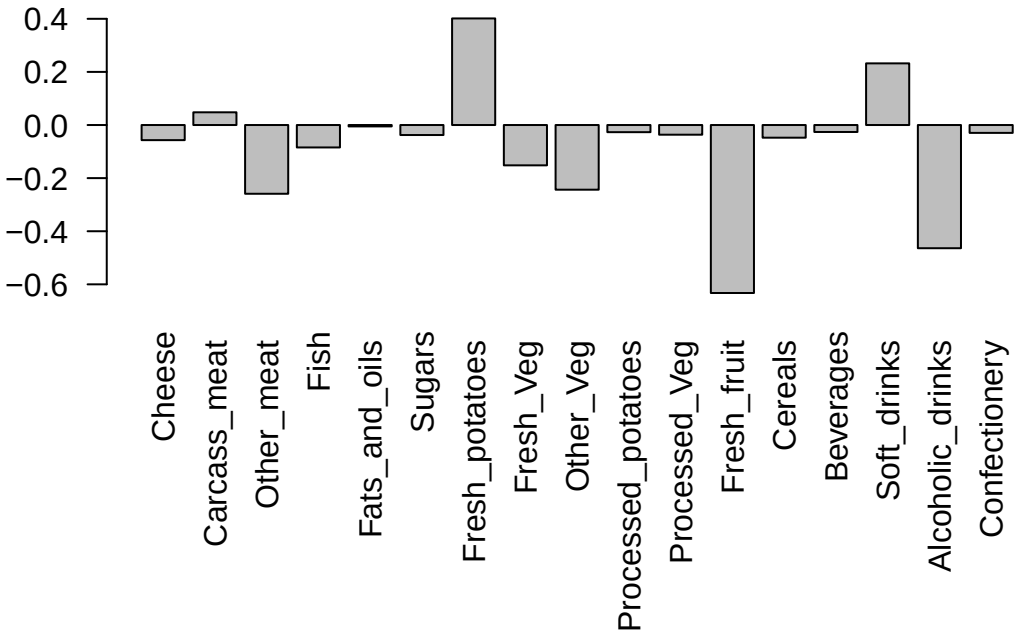

Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominantely and what does PC2 maninly tell us about?

Answer: Fresh potatoes and soft drinks are the two food groups that are featured prominately in the 'loadings plot' for PC2. PC2 details the second-largest variance in the dataset, which means it captures variations not included in PC1.

```
## Lets focus on PC2
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```

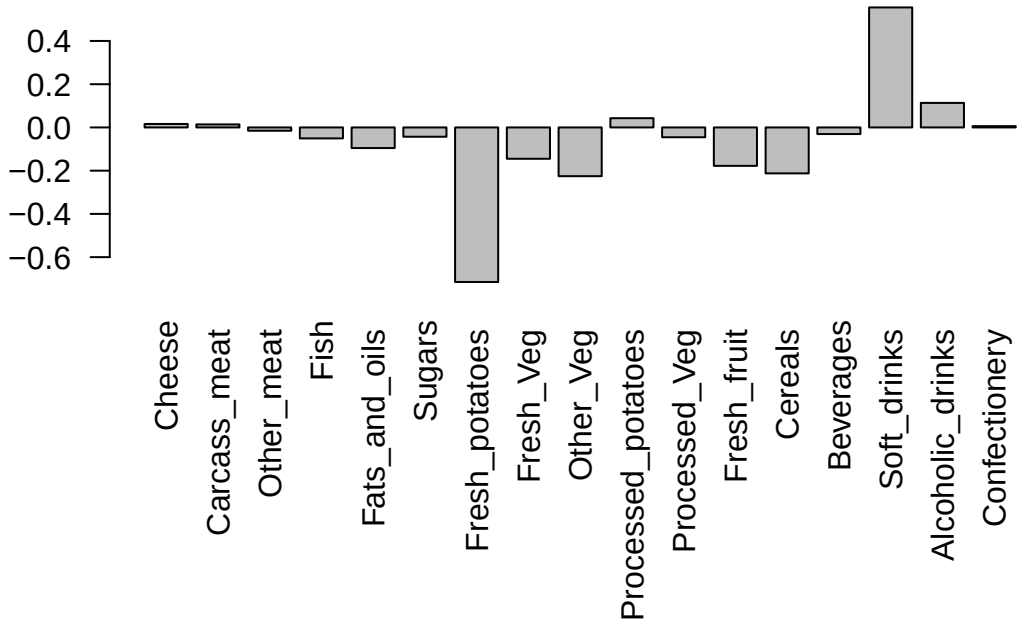

Using ggplot for the figures

```
library(ggplot2)

df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```

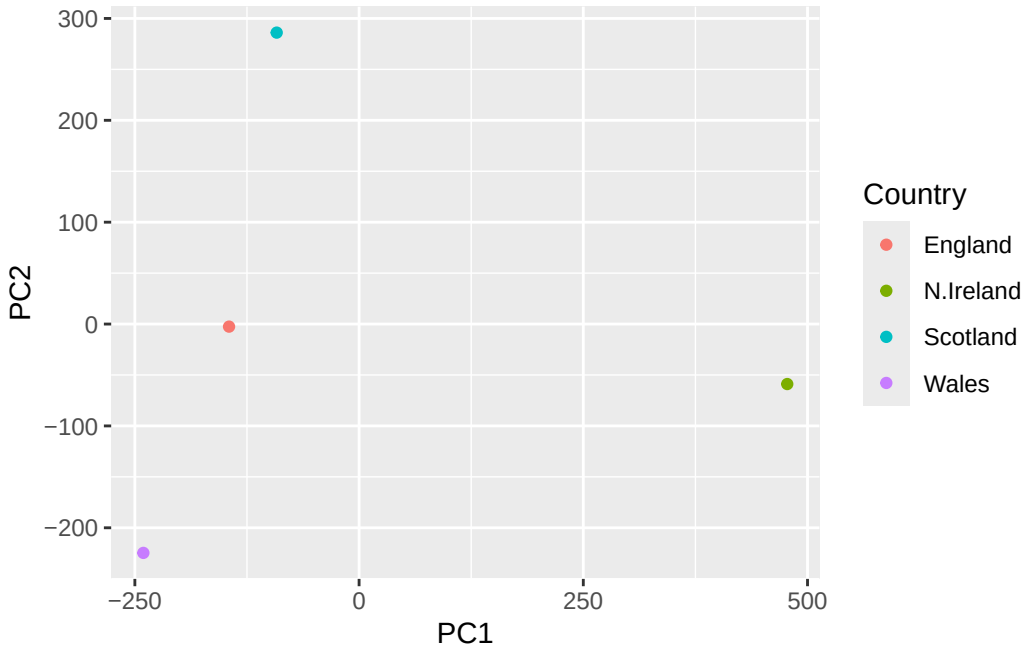


```
# a much nicer plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country, label=Country) +
  geom_hline(yintercept = 0, col="gray") +
  geom_vline(xintercept = 0, col="gray") +
  geom_point(show.legend = FALSE) +
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +
  expand_limits(x = c(-300,500)) +
  xlab("PC1 (67.4%)") +
  ylab("PC2 (28%)") +
  theme_bw()
```

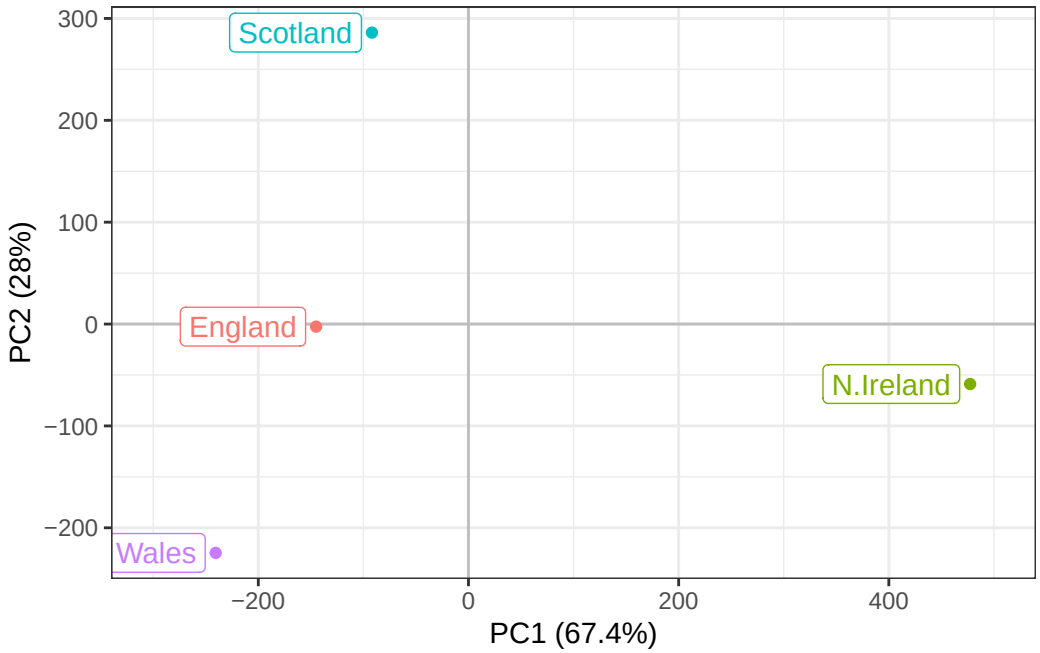


```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")

ggplot(ld_lab) +
  aes(PC1, Food) +
  geom_col()
```

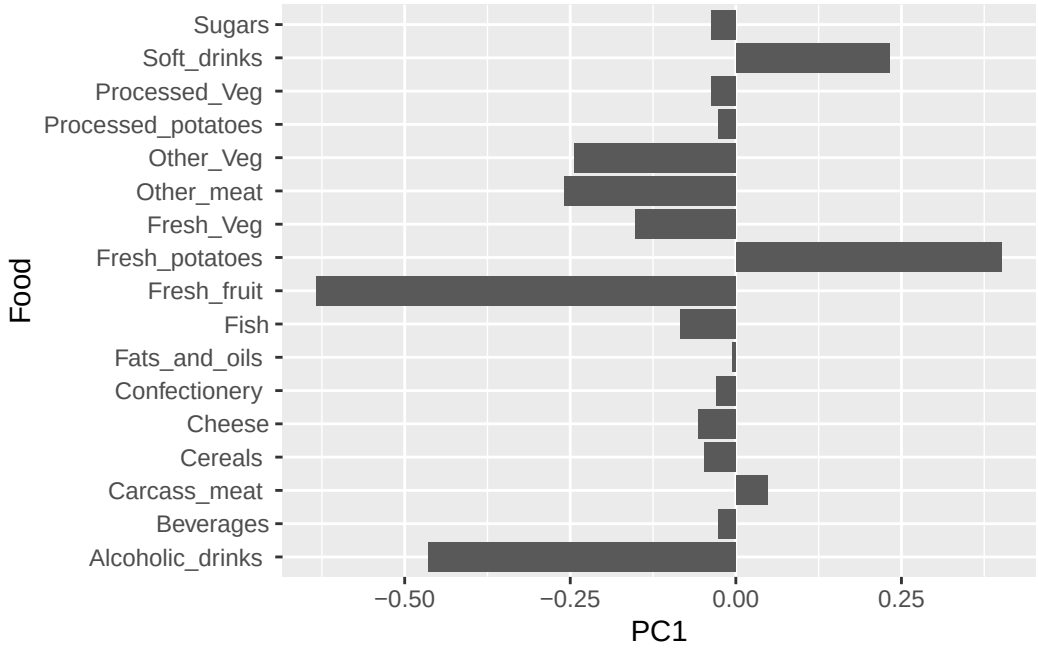


```
## PCA of RNA-seq data
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
       wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
gene1  439 458  408  429 420  90  88  86  90  93
gene2  219 200  204  210 187 427 423 434 433 426
gene3 1006 989 1030 1017 973 252 237 238 226 210
gene4  783 792  829  856 760 849 856 835 885 894
gene5  181 249  204  244 225 277 305 272 270 279
gene6  460 502  491  491 493 612 594 577 618 638
```

Q10: How many genes and samples are in this data set?

Answer:There are 100 rows, which means there is 100 genes. There are 10 columns, which means there are 10 samples.

```
nrow(rna.data)
```

```
[1] 100
```

```
ncol(rna.data)
```

```
[1] 10
```

```
dim(rna.data)
```

```
[1] 100  10
```